

Protect Before Losing: Data-Driven Early Warning System for Enhancing Student Retention in Higher Education

Literature Review, Data Research, and Technology Review

Team Members:

Şiyar Kafuroğulları, Yaren Yeşiltepe, Murat Ertik, Mehmet Faruk Bağdaş

SECTION 1: LITERATURE REVIEW

1. Introduction

In higher education, a student is not merely an academic participant but a key customer and stakeholder within the institutional ecosystem. Across OECD countries, approximately 30% of students drop out before completing their higher education. From a marketing perspective, student dropout represents a critical **Customer Retention and Churn Problem**. When a student leaves an institution prematurely, it results in lost lifetime value (LTV), wasted recruitment and scholarship investment, and damage to institutional brand equity and university rankings.

This project addresses the challenge of identifying students at risk of churning during their first academic year using demographic, macroeconomic, and academic metrics.

Primary Research Question: How can machine learning models be leveraged as a predictive Customer Relationship Management (CRM) tool to identify early churn signals and improve student retention rate (KPI) in higher education?

2. Organization of Literature

The existing literature is organized around two main themes:

- Predictive Churn Modeling & Student Retention in Higher Education:** Studies applying machine learning to identify early indicators of student attrition.
- CRM and Proactive Engagement Strategies in Service Operations:** Literature linking early warning systems to customer success and targeted retention interventions.

3. Summary and Synthesis of Literature

- **Realinho et al. (2022) — Predicting Student Dropout and Academic Success:** This study established a benchmark dataset (UCI Repository) combining demographic, macroeconomic, and academic performance attributes. The authors applied tree-based ensemble algorithms (Random Forest, Gradient Boosting) to classify students into three categories: Dropout, Enrolled, and Graduate, achieving high predictive accuracy. Their

findings demonstrated that academic performance indicators in the first semester are the strongest predictors of dropout.

- **Aulck et al. (2017) — *Predicting Student Dropout in Higher Education*:** Analyzing large-scale historical transcript data, Aulck et al. implemented Random Forest, Logistic Regression, and k-NN models to predict first-year dropout. They highlighted that GPA, course load, and early academic friction serve as primary early warning indicators, proving that early intervention is significantly more cost-effective than late-stage re-engagement.

Synthesis: Both studies prove that structured tabular data containing early academic and background indicators provides sufficient signal for predicting churn. However, while existing literature predominantly treats dropout prediction as a purely administrative or pedagogical issue, there is a gap in translating these predictive outputs into operationalized **Marketing/CRM Workflows** for targeted student engagement and brand loyalty management.

4. Marketing Relevance and Gap

Current academic approaches successfully identify *who* will drop out, but often fail to integrate these insights into actionable marketing and customer success pipelines. Academic advisory decisions are often driven by assumptions rather than data-driven Key Performance Indicators (KPIs). This capstone project bridges this gap by positioning the predictive model as an early warning CRM engine that enables personalized, targeted retention strategies before the first academic year concludes, directly improving institutional **Student Retention Rates** and **Brand Reputation**.

5. Conclusion

The literature confirms that tree-based machine learning models trained on early-stage student data yield high predictive accuracy for churn. Integrating these models into a university's CRM framework allows for proactive, evidence-based intervention, optimizing resource allocation and boosting student lifetime loyalty.

SECTION 2: DATA RESEARCH

1. Introduction

To build a reliable churn prediction engine, the data must capture multi-dimensional aspects of the customer (student) journey, including pre-enrollment background, socio-economic context, and early post-enrollment performance. The selected dataset enables the predictive model to output a risk score for each student, serving as a core marketing KPI for academic advisors.

2. Data Source and Scope

- **Data Source:** UCI Machine Learning Repository — "*Predict Students' Dropout and Academic Success*" dataset.
- **Data Type:** Public, structured tabular data (.csv).
- **Scale:** 4,424 student records and 36 attributes.
- **Features:**
 - **Demographic & Socioeconomic:** Age at enrollment, gender, marital status, nationality, parents' qualification/occupation, scholarship holder, debtor status, tuition fees up to date.

- **Macroeconomic:** Unemployment rate, inflation rate, GDP.
- **Academic Performance (1st & 2nd Semester):** Enrolled units, approved units, evaluations, grade averages, application mode/order.
- **Target Variable:** Student status classified into three categories: *Dropout*, *Enrolled*, and *Graduate*.

3. Data Quality, Privacy, and Limitations

- **Class Imbalance:** The target variable exhibits class imbalance across Dropout, Enrolled, and Graduate outcomes. To address this, **SMOTE (Synthetic Minority Over-sampling Technique)** will be applied during preprocessing.
- **Preprocessing Pipeline:**
 1. *Encoding:* One-Hot Encoding for categorical variables and Label Encoding for ordinal/target categories.
 2. *Scaling:* StandardScaler applied to continuous variables (age, GPA, inflation rate).
 3. *Splitting:* Stratified 80/20 train-test split to preserve class distribution.
- **Privacy & Ethics:** The dataset is anonymized and publicly benchmarked, compliant with personal data protection guidelines.

4. Exploratory Data Analysis & Marketing Insights

Preliminary EDA highlights key relationships relevant to customer retention KPIs:

- **Tuition Fee Up-to-Date & Debtor Status:** Students with overdue tuition or unpaid debts exhibit a significantly higher churn risk, mirroring financial friction points in commercial subscription models.
- **1st Semester Approved Units & GPA:** Early academic failure (low approved units) is the single strongest driver of early churn.
- **Scholarship Support:** Scholarship recipients demonstrate markedly higher retention rates, underlining the impact of targeted financial incentives on customer loyalty.

5. Conclusion

The UCI dataset provides a robust, clean, and representative foundation for training supervised classification models. The attributes directly map to actionable marketing levers (e.g., financial aid targeting, specialized advisory campaigns) necessary for the next project phase.

SECTION 3: TECHNOLOGY REVIEW

1. Introduction & Technology Overview

The technical solution centers on **Tree-Based Ensemble Machine Learning Algorithms**, specifically **Random Forest**, **XGBoost**, and **LightGBM**. In structured, medium-sized tabular datasets, tree ensembles consistently outperform deep learning approaches while offering vastly superior model interpretability.

2. Relevance to Marketing Project & Workflow

In a modern marketing tech stack, predictive models power automated customer segmentation and personalized communication. The technology fits into the university's operational CRM as follows:

- 1. Data Ingestion:** Student enrollment and 1st-term performance data are fed into the pipeline.
- 2. Predictive Scoring:** The XGBoost/LightGBM model evaluates risk probabilities before the completion of the first year.
- 3. Automated CRM Segmentation:** High-risk students are tagged in the CRM dashboard.
- 4. Targeted Campaign Execution:** Academic advisors receive automated early warnings to execute tailored retention interventions (e.g., 1-on-1 counseling, tutoring, financial aid consultation).

3. Technology Comparison and Evaluation

Evaluation Criteria	Random Forest	XGBoost / LightGBM	Neural Networks (Deep Learning)
Tabular Data Performance	High	Superior (State-of-the-Art)	Moderate / Requires massive data
Overfitting Resistance	High	High (with regularization)	Low on small/medium datasets
Interpretability (SHAP)	Excellent	Excellent	Low ("Black Box")
Training Speed & Scalability	Fast	Extremely Fast	Computationally Heavy

4. Model Explainability (SHAP Values)

In higher education marketing, explaining *why* a student is flagged as high-risk is as crucial as the prediction itself. By computing **SHAP (SHapley Additive exPlanations)** values, advisors can inspect individual risk factors (e.g., financial strain vs. academic difficulty) and deploy the precise intervention needed, ensuring evidence-based decision-making.

5. Real-World Marketing Use Cases

Similar tree-based churn prediction architectures are widely deployed in commercial SaaS, telecom, and financial service industries to trigger proactive customer success workflows, resulting in reduced churn rates and increased customer lifetime value (CLV).

6. Limitations, Risks, and Mitigation

- Data Drift:** Student demographics or grading policies may change over time. *Mitigation:* Periodic model retraining with fresh cohort data.
- False Positives:** Flagging a low-risk student may cause unnecessary outreach. *Mitigation:* Human-in-the-loop validation by human academic advisors before outreach.

7. Conclusion

Tree-based ensemble models (Random Forest / XGBoost / LightGBM) represent the optimal technological choice for this capstone. They offer top-tier predictive accuracy on tabular data, low latency, and transparent model explainability (SHAP), providing quantifiable business value and enhancing student retention.

REFERENCES

1. Aulck, L., Velagapudi, N., Blumenstock, J., & West, J. (2017). Predicting student dropout in higher education. *arXiv preprint arXiv:1707.04309*.
2. Realinho, V., Machado, J., Baptista, A., & Martins, M. V. (2022). Predicting student dropout and academic success. *Data*, 7(11), 146. UCI Machine Learning Repository.
3. Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems (NeurIPS)*, 30.